

Simran Kaur

311443

DDA Lab 7 Report

In [1]:

```
import torch
from torch.nn import Module
from torch import nn
import torchvision
import torchvision.transforms as transforms
from torch.utils.data import DataLoader, Dataset
from torch.utils.tensorboard import SummaryWriter
import torch.nn.functional as F
import matplotlib.pyplot as plt
import numpy as np
```

Network Analysis: Image Classification - Part 2

The Network has three convolution layer, two max pool and three fully connected layers with ReLU as activation function. Since the Loss here used is CrossEntropy so there wasn't any need for the softmax layer. The input and output channels are selected after some experiments.

In [2]:

```
class CNNetwork(Module):
    def __init__(self, input_channels):
        super(CNNetwork, self).__init__()
        self.conv1 = nn.Conv2d(input_channels, 10, kernel_size = 5)
        self.pool1 = nn.MaxPool2d(kernel_size = 2, stride = 2)
        self.conv2 = nn.Conv2d(10, 30, kernel_size = 5)
        self.conv3 = nn.Conv2d(30, 60, kernel_size = 5)
        self.pool2 = nn.MaxPool2d(kernel_size = 2, stride = 2)
        self.fc1 = nn.Linear(540, 270)
        self.fc2 = nn.Linear(270, 130)
        self.fc3 = nn.Linear(130, 10)

    def forward(self, x):
        y = F.relu(self.conv1(x))
        y = self.pool1(y)
        y = F.relu(self.conv2(y))
        y = F.relu(self.conv3(y))
        y = self.pool2(y)
        y = y.view(-1, 60*3*3)
        y = F.relu(self.fc1(y))
        y = F.relu(self.fc2(y))
        y = self.fc3(y)
        return y
```

In [3]:

```
CifarTrain = torchvision.datasets.CIFAR10(root = './data', train = True, transform = transf
CifarTest = torchvision.datasets.CIFAR10(root = './data', train = False, transform = transf
```

Files already downloaded and verified
Files already downloaded and verified

In [6]:

```
batch_size = 128
epochs = 10
learning_rate = 10**(-3)
CifarTrainHalf = torch.utils.data.Subset(CifarTrain, range(len(CifarTrain)//2))
```

In [7]:

```
train_loader = DataLoader(CifarTrainHalf, batch_size = batch_size, shuffle=True)
test_loader = DataLoader(CifarTest, batch_size = 64, shuffle = True)
```

In []:

```

ModelCNN = CNNNetwork(3)
criterion = nn.CrossEntropyLoss()
optimizer = torch.optim.Adam(ModelCNN.parameters(), lr = learning_rate)
writer = SummaryWriter(f'runs/Cifar/Network')
step = 0
for epoch in range(epochs):
    total_loss = 0
    correct_pred = 0
    sample_size = 0
    for i, batch in enumerate(train_loader):
        optimizer.zero_grad()
        yhat = ModelCNN(batch[0])
        loss = criterion(yhat.squeeze(), batch[1].squeeze())
        total_loss += loss.item()
        pred = torch.max(yhat.data, 1)[1]
        sample_size = batch[1].size(0)
        correct_pred += (pred == batch[1]).sum().item()
        loss.backward()
        optimizer.step()
        if (i + 1) % 50 == 0:
            writer.add_scalar('Loss on Minibatches', loss.item(), step + 1)
            writer.add_scalar('Accuracy on minibatches', ((pred == batch[1]).sum().item() / sample_size),
                               step + 1)
            step += 1
    accuracy = (correct_pred / len(CifarTrainHalf)) * 100
    writer.add_scalar('Train Loss', np.sqrt(total_loss / len(CifarTrainHalf)), epoch + 1)
    writer.add_scalar('Train accuracy', accuracy, epoch + 1)


    with torch.no_grad():
        test_pred = 0
        test_loss = 0
        for i, (img, label) in enumerate(test_loader):
            optimizer.zero_grad()
            output = ModelCNN(img)
            loss = criterion(output.squeeze(), label.squeeze())
            test_loss += loss.item()
            pred = torch.max(output.data, 1)[1]
            test_pred += (pred == label).sum().item()
        accuracy = (test_pred / len(CifarTest)) * 100
        writer.add_scalar('Test Loss', np.sqrt(test_loss / len(CifarTest)), epoch + 1)
        writer.add_scalar('Test accuracy', accuracy, epoch + 1)


writer.close()
writer.flush()

```

For batches we can see the trend in loss and accuracy, as the model sees more and more data, loss is decreased and accuracy is increased.


Loss on Minibatches



Name	Smoothed	Value	Step	Time	Relative
 Cifar\Network	0.974	0.974	30	Sun Jun 26, 17:10:37	1m 58s

Accuracy on minibatches





Name	Smoothed	Value	Step	Time	Relative
 Cifar\Network	67.97	67.97	30	Sun Jun 26, 17:10:37	1m 58s

Training and Test Loss


For the 10th epoch, the test loss as seen below is very high (0.14) as compared to the training loss which is 0.09


Train Loss



Name	Smoothed	Value	Step	Time	Relative
 Cifar\Network	0.0924	0.0924	10	Sun Jun 26, 17:10:40	1m 52s

Test Loss




Name	Smoothed	Value	Step	Time	Relative
 Cifar\Network	0.1387	0.1387	10	Sun Jun 26, 17:10:42	1m 52s

Training and Test Accuracy


Also on comparing accuracy we see that training accuracy is more than the test accuracy

Train accuracy



	Name	Smoothed	Value	Step	Time	Relative
	Cifar\Network	61.23	61.23	10	Sun Jun 26, 17:10:40	1m 52s

Test accuracy



	Name	Smoothed	Value	Step	Time	Relative
	Cifar\Network	55.9	55.9	10	Sun Jun 26, 17:10:42	1m 52s

Now we will see approaches to deal with overfitting as seen above

Normalization Effect

mean and standard deviation for Normalization

In [8]:

```
R_total_mean, G_total_mean, B_total_mean = 0, 0, 0
R_total_std, G_total_std, B_total_std = 0, 0, 0
for img, label in CifarTrain:
    R_mean, G_mean, B_mean = torch.mean(img, dim = [1,2])
    R_std, G_std, B_std = torch.std(img, dim = [1,2])
    R_total_mean += R_mean
    G_total_mean += G_mean
    B_total_mean += B_mean
    R_total_std += R_std
    G_total_std += G_std
    B_total_std += B_std
print(R_total_mean/len(CifarTrain), G_total_mean/len(CifarTrain), B_total_mean/len(CifarTrain))
print(R_total_std/len(CifarTrain), G_total_std/len(CifarTrain), B_total_std/len(CifarTrain))
```

```
tensor(0.4914) tensor(0.4822) tensor(0.4465)
tensor(0.2023) tensor(0.1994) tensor(0.2010)
```

Transformations defined for augmentation, normalization and both as below

In []:

```
transformAug = transforms.Compose([transforms.Resize((32,32)),
                                   transforms.RandomHorizontalFlip(),
                                   transforms.RandomRotation(10),
                                   transforms.ColorJitter(brightness=0.3, contrast=0.2,
                                                         saturation=0.1, hue=0.1),
                                   transforms.ToTensor()
                                   ])

transformNorm = transforms.Compose([transforms.Resize((32,32)),
                                   transforms.ToTensor(),
                                   transforms.Normalize((0.4914, 0.4822, 0.4465), (0.2023, 0.1994, 0.2010))
                                   ])

transformBoth = transforms.Compose([transforms.Resize((32,32)),
                                   transforms.RandomHorizontalFlip(),
                                   transforms.RandomRotation(10),
                                   transforms.ColorJitter(brightness=0.3, contrast=0.2,
                                                         saturation=0.1, hue=0.1),
                                   transforms.ToTensor(),
                                   transforms.Normalize((0.4914, 0.4822, 0.4465), (0.2023, 0.1994, 0.2010))
                                   ])
```

1. Augmentation

The general idea is to increase dataset by appending more images with some changes such as change in contrast, rotation etc.

In []:

```
CifarTrainAug = torchvision.datasets.CIFAR10(root = './data', train = True,  
                                             transform = transformAug, download = True)  
CifarTestAug = torchvision.datasets.CIFAR10(root = './data', train = False,  
                                             transform = transformAug, download = True)
```

In []:

```
AugmentedDataTrain = torch.utils.data.ConcatDataset([CifarTrain, CIFarTrainAug])  
AugmentedDataTest = torch.utils.data.ConcatDataset([CifarTest, CIFarTestAug])
```

In []:

```
train_loaderAug = DataLoader(AugmentedDataTrain, batch_size = batch_size, shuffle=True)  
test_loaderAug = DataLoader(AugmentedDataTest, batch_size = 64, shuffle = True)
```


In []:

```


ModelCNNAug = CNNNetwork(3)
criterion = nn.CrossEntropyLoss()
optimizer = torch.optim.Adam(ModelCNNAug.parameters(), lr = learning_rate)
writer = SummaryWriter(f'runs/Cifar/Augmentation')
step = 0
for epoch in range(epochs):
    total_loss = 0
    correct_pred = 0
    sample_size = 0
    for i, batch in enumerate(train_loaderAug):
        optimizer.zero_grad()
        yhat = ModelCNNAug(batch[0])
        loss = criterion(yhat.squeeze(), batch[1].squeeze())
        total_loss += loss.item()
        pred = torch.max(yhat.data, 1)[1]
        sample_size = batch[1].size(0)
        correct_pred += (pred == batch[1]).sum().item()
        loss.backward()
        optimizer.step()
        if (i + 1) % 50 == 0:
            writer.add_scalar('Loss on Minibatches', loss.item(), step + 1)
            writer.add_scalar('Accuracy on minibatches', ((pred == batch[1]).sum().item() / sample_size),
                              step + 1)
            step += 1
    accuracy = (correct_pred / len(AugmentedDataTrain)) * 100
    writer.add_scalar('Train Loss', np.sqrt(total_loss / len(AugmentedDataTrain)), epoch + 1)
    writer.add_scalar('Train accuracy', accuracy, epoch + 1)

    with torch.no_grad():
        test_pred = 0
        test_loss = 0
        for i, (img, label) in enumerate(test_loaderAug):
            optimizer.zero_grad()
            output = ModelCNNAug(img)
            loss = criterion(output.squeeze(), label.squeeze())
            test_loss += loss.item()
            pred = torch.max(output.data, 1)[1]
            test_pred += (pred == label).sum().item()
        accuracy = (test_pred / len(AugmentedDataTest)) * 100
        writer.add_scalar('Test Loss', np.sqrt(test_loss / len(AugmentedDataTest)), epoch + 1)
        writer.add_scalar('Test accuracy', accuracy, epoch + 1)

writer.close()
writer.flush()


```

Loss on Minibatches



Name	Smoothed	Value	Step	Time	Relative
Cifar\Augmentation	0.76	0.76	150	Sun Jun 26, 17:32:34	11m 46s

Accuracy on minibatches




Name	Smoothed	Value	Step	Time	Relative
Cifar\Augmentation	74.22	74.22	150	Sun Jun 26, 17:32:34	11m 46s

Training and Test Loss

Comparison with baseline


The Network seems to work even better after Augmentation to the dataset. Not only the loss on training has decreased but also on the test itself.

Train Loss



Name	Smoothed	Value	Step	Time	Relative
Cifar\Augmentation	0.07862	0.07862	10	Sun Jun 26, 17:32:37	10m 48s
Cifar\Network	0.0924	0.0924	10	Sun Jun 26, 17:10:40	1m 52s

Test Loss



Name	Smoothed	Value	Step	Time	Relative
Cifar\Augmentation	0.1248	0.1248	10	Sun Jun 26, 17:32:46	10m 48s
Cifar\Network	0.1387	0.1387	10	Sun Jun 26, 17:10:42	1m 52s

Training and test accuracy



comparison with baseline

After augmentation, the training and test accuracy have both risen.

Train Loss

	Name	Smoothed	Value	Step	Time	Relative
Train accuracy	<div><div></div>Cifar\Augmentation</div>	72.18	72.18	10	Sun Jun 26, 17:32:37	10m 48s
	<div><div></div>Cifar\Network</div>	61.23	61.23	10	Sun Jun 26, 17:10:40	1m 52s

Train accuracy



	Name	Smoothed	Value	Step	Time	Relative
Train Loss	<div><div></div>Cifar\Augmentation</div>	65.93	65.93	10	Sun Jun 26, 17:32:46	10m 48s
	<div><div></div>Cifar\Network</div>	55.9	55.9	10	Sun Jun 26, 17:10:42	1m 52s

Normalization

Normalizing dataset also has huge impact on the performance of the model.

In []:

```
CifarTrainNorm = torchvision.datasets.CIFAR10(root = './data', train = True,  
                                              transform = transformNorm, download = True)  
CifarTestNorm = torchvision.datasets.CIFAR10(root = './data', train = False,  
                                              transform = transformNorm, download = True)
```

In []:

```
train_loaderNorm = DataLoader(CifarTrainNorm, batch_size = batch_size, shuffle=True)  
test_loaderNorm = DataLoader(CifarTestNorm, batch_size = 64, shuffle = True)
```

In []:

```

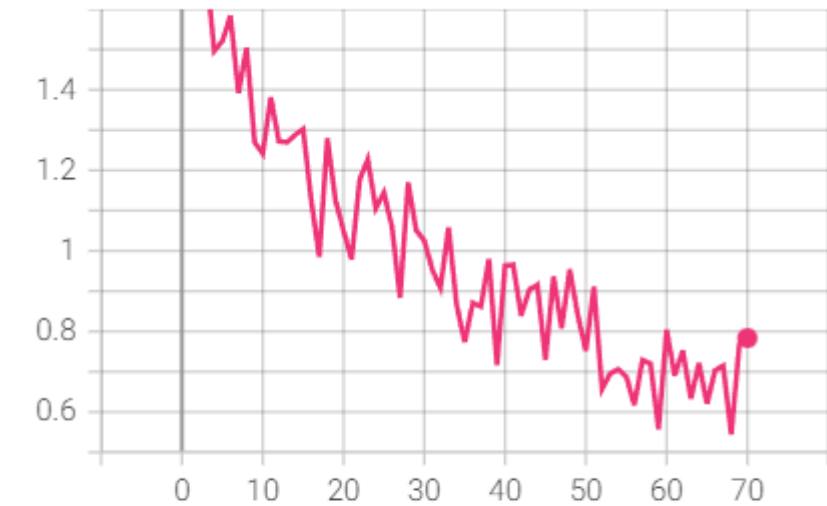
ModelCNNnorm = CNNNetwork(3)
criterion = nn.CrossEntropyLoss()
optimizer = torch.optim.Adam(ModelCNNnorm.parameters(),lr = learning_rate)
writer = SummaryWriter(f'runs/Cifar/Normalization')
step = 0
for epoch in range(epochs):
    total_loss = 0
    correct_pred = 0
    sample_size = 0
    for i, batch in enumerate(train_loaderNorm):
        optimizer.zero_grad()
        yhat = ModelCNNnorm(batch[0])
        loss = criterion(yhat.squeeze(), batch[1].squeeze())
        total_loss += loss.item()
        pred = torch.max(yhat.data, 1)[1]
        sample_size = batch[1].size(0)
        correct_pred += (pred == batch[1]).sum().item()
        loss.backward()
        optimizer.step()
        if (i + 1) % 50 == 0:
            writer.add_scalar('Loss on Minibatches', loss.item(), step + 1)
            writer.add_scalar('Accuracy on minibatches',((pred == batch[1]).sum().item()/sa
                step + 1)
            step += 1
    accuracy = (correct_pred/len(CifarTrainNorm))*100
    writer.add_scalar('Train Loss', np.sqrt(total_loss/len(CifarTrainNorm)), epoch + 1)
    writer.add_scalar('Train accuracy', accuracy, epoch + 1)




    with torch.no_grad():
        test_pred = 0
        test_loss = 0
        for i, (img, label) in enumerate(test_loaderNorm):
            optimizer.zero_grad()
            output = ModelCNNnorm(img)
            loss = criterion(output.squeeze(), label.squeeze())
            test_loss += loss.item()
            pred = torch.max(output.data, 1)[1]
            test_pred += (pred == label).sum().item()
        accuracy = (test_pred/len(CifarTestNorm))*100
        writer.add_scalar('Test Loss', np.sqrt(test_loss/len(CifarTestNorm)), epoch + 1)
        writer.add_scalar('Test accuracy', accuracy, epoch + 1)


writer.close()
writer.flush()

```


Loss on Minibatches









Name	Smoothed	Value	Step	Time	Relative
 Cifar\Normalization	0.7836	0.7836	70	Sun Jun 26, 17:42:24	4m 47s

Accuracy on minibatches






Name	Smoothed	Value	Step	Time	Relative
 Cifar\Normalization	69.53	69.53	70	Sun Jun 26, 17:42:24	4m 47s

Training and Test Loss

Comparison with baseline


After normalization both training and test loss has decreased as can be seen below.

Train Loss



Name	Smoothed	Value	Step	Time	Relative
Cifar\Network	0.0924	0.0924	10	Sun Jun 26, 17:10:40	1m 52s
Cifar\Normalization	0.07022	0.07022	10	Sun Jun 26, 17:42:27	4m 26s

Test Loss




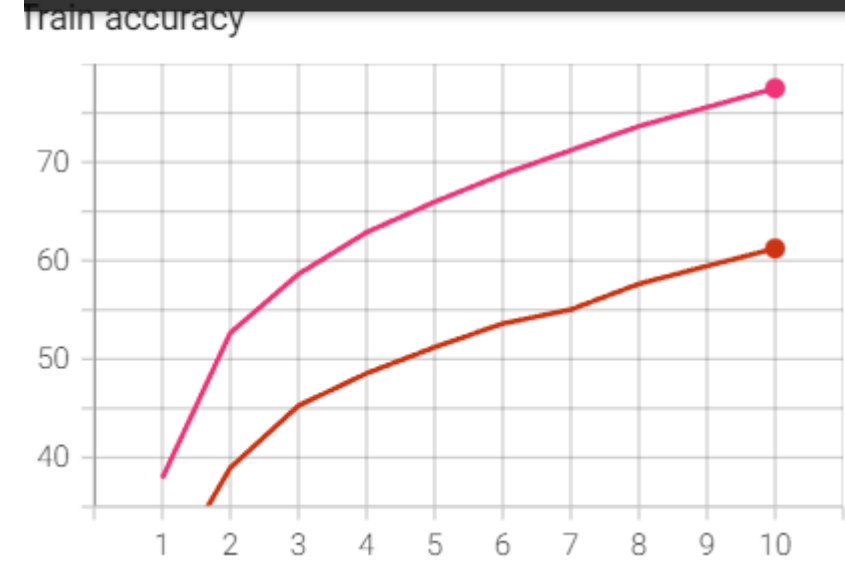
Name	Smoothed	Value	Step	Time	Relative
Cifar\Network	0.1387	0.1387	10	Sun Jun 26, 17:10:42	1m 52s
Cifar\Normalization	0.1295	0.1295	10	Sun Jun 26, 17:42:30	4m 26s

Training and test accuracy comparison with baseline

The training accuracy has increased from around 60% to 77.5% while the test accuracy from 60% to 66% over a range of 10 epochs.

Train Loss

	Name	Smoothed	Value	Step	Time	Relative
●	Cifar\Network	61.23	61.23	10	Sun Jun 26, 17:10:40	1m 52s
●	Cifar\Normalization	77.5	77.5	10	Sun Jun 26, 17:42:27	4m 26s



☐ ☐ ☐

	Name	Smoothed	Value	Step	Time	Relative
●	Cifar\Network	55.9	55.9	10	Sun Jun 26, 17:10:42	1m 52s
●	Cifar\Normalization	66.24	66.24	10	Sun Jun 26, 17:42:30	4m 26s

Augmentation + Normalization

Now both the Augmentation and normalization has been applied to the dataset.

In []:

```
CifarTrainAugNorm = torchvision.datasets.CIFAR10(root = './data', train = True,
                                                  transform = transformBoth, download = True)
CifarTestAugNorm = torchvision.datasets.CIFAR10(root = './data', train = False,
                                                  transform = transformBoth, download = True)
```

In []:

```
AugmentedTrain = torch.utils.data.ConcatDataset([CifarTrainAugNorm, CifarTrainNorm])
AugmentedTest = torch.utils.data.ConcatDataset([CifarTestAugNorm, CifarTestNorm])
```

In []:

```
train_loaderAugNorm = DataLoader(AugmentedTrain, batch_size = batch_size, shuffle=True)
test_loaderAugNorm = DataLoader(AugmentedTest, batch_size = 64, shuffle = True)
```

In []:

```


Model = CNNNetwork(3)
criterion = nn.CrossEntropyLoss()
optimizer = torch.optim.Adam(Model.parameters(), lr = learning_rate)
writer = SummaryWriter(f'runs/Cifar/Norm_Aug')
step = 0
for epoch in range(epochs):
    total_loss = 0
    correct_pred = 0
    sample_size = 0
    for i, batch in enumerate(train_loaderAugNorm):
        optimizer.zero_grad()
        yhat = Model(batch[0])
        loss = criterion(yhat.squeeze(), batch[1].squeeze())
        total_loss += loss.item()
        pred = torch.max(yhat.data, 1)[1]
        sample_size = batch[1].size(0)
        correct_pred += (pred == batch[1]).sum().item()
        loss.backward()
        optimizer.step()
        if (i + 1) % 50 == 0:
            writer.add_scalar('Loss on Minibatches', loss.item(), step + 1)
            writer.add_scalar('Accuracy on minibatches', ((pred == batch[1]).sum().item() / sample_size),
                              step + 1)
            step += 1
    accuracy = (correct_pred / len(AugmentedTrain)) * 100
    writer.add_scalar('Train Loss', np.sqrt(total_loss / len(AugmentedTrain)), epoch + 1)
    writer.add_scalar('Train accuracy', accuracy, epoch + 1)

    with torch.no_grad():
        test_pred = 0
        test_loss = 0
        for i, (img, label) in enumerate(test_loaderAugNorm):
            optimizer.zero_grad()
            output = Model(img)
            loss = criterion(output.squeeze(), label.squeeze())
            test_loss += loss.item()
            pred = torch.max(output.data, 1)[1]
            test_pred += (pred == label).sum().item()
        accuracy = (test_pred / len(AugmentedTest)) * 100
        writer.add_scalar('Test Loss', np.sqrt(test_loss / len(AugmentedTest)), epoch + 1)
        writer.add_scalar('Test accuracy', accuracy, epoch + 1)

writer.close()
writer.flush()


```

Loss on Minibatches



	Name	Smoothed	Value	Step	Time	Relative
Test Loss	Cifar\Norm_Aug	0.6916	0.6916	150	Sun Jun 26, 15:08:31	13m 25s

Accuracy on minibatches




	Name	Smoothed	Value	Step	Time	Relative
Loss on minibatches	Cifar\Norm_Aug	78.13	78.13	150	Sun Jun 26, 15:08:31	13m 25s

Training and Test Loss


Comparison with baseline

Train Loss



	Name	Smoothed	Value	Step	Time	Relative
Train accuracy	Cifar\Network	0.0924	0.0924	10	Sun Jun 26, 17:10:40	1m 52s
	Cifar\Norm_Aug	0.06956	0.06956	10	Sun Jun 26, 15:08:34	12m 21s

Test Loss



	Name	Smoothed	Value	Step	Time	Relative
Test accuracy	Cifar\Network	0.1387	0.1387	10	Sun Jun 26, 17:10:42	1m 52s
	Cifar\Norm_Aug	0.1215	0.1215	10	Sun Jun 26, 15:08:44	12m 21s



Training and Test accuracy

Comparison with baseline

This has resulted in the higher accuracy increase so far with train accuracy around 79% and test accuracy 70% over a range of 10 epochs.

Train Loss

	Name	Smoothed	Value	Step	Time	Relative
●	Cifar\Network	61.23	61.23	10	Sun Jun 26, 17:10:40	1m 52s
○	Cifar\Norm_Aug	78.43	78.43	10	Sun Jun 26, 15:08:34	12m 21s



Train Loss

	Name	Smoothed	Value	Step	Time	Relative
●	Cifar\Network	55.9	55.9	10	Sun Jun 26, 17:10:42	1m 52s
○	Cifar\Norm_Aug	69.32	69.32	10	Sun Jun 26, 15:08:44	12m 21s

Network Regularization

So far the modification to the dataset has resulted in the increase in the model accuracy as well as decrease in the loss. These all modification changes from model to model as well as the amount of training data available. Now we will see how can we make changes in the model itself, by dropping some neurons or by penalising the weights associated with them to improve model efficiency.

Dropout

For the fully connected layers, we give a probability for the number of neurons that can be dropped that the model over training finds to be not that important in order to reduce complexity as well as increase efficiency.

In []:

```
class CNNDropout(Module):
    def __init__(self, input_channels, dropProb1, dropProb2):
        super(CNNNetwork, self).__init__()
        self.conv1 = nn.Conv2d(input_channels, 10, kernel_size = 5)
        self.pool1 = nn.MaxPool2d(kernel_size = 2, stride = 2)
        self.conv2 = nn.Conv2d(10, 30, kernel_size = 5)
        self.conv3 = nn.Conv2d(30, 60, kernel_size = 5)
        self.pool2 = nn.MaxPool2d(kernel_size = 2, stride = 2)
        self.fc1 = nn.Linear(540, 270)
        self.dropout1 = nn.Dropout(dropProb1)
        self.fc2 = nn.Linear(270, 130)
        self.dropout2 = nn.Dropout(dropProb2)
        self.fc3 = nn.Linear(130, 10)

    def forward(self, x):
        y = F.relu(self.conv1(x))
        y = self.pool1(y)
        y = F.relu(self.conv2(y))
        y = F.relu(self.conv3(y))
        y = self.pool2(y)
        y = y.view(-1, 60*3*3)
        y = F.relu(self.fc1(y))
        y = self.dropout1(y)
        y = F.relu(self.fc2(y))
        y = self.dropout2(y)
        y = self.fc3(y)
        return y
```

In []:

```


ModelD = CNNDropout(3, 0.3, 0.2)
criterion = nn.CrossEntropyLoss()
optimizer = torch.optim.Adam(ModelD.parameters(), lr = learning_rate)
writer = SummaryWriter(f'runs/Cifar/Dropout')
step = 0
for epoch in range(epochs):
    total_loss = 0
    correct_pred = 0
    sample_size = 0
    for i, batch in enumerate(train_loader):
        optimizer.zero_grad()
        yhat = ModelD(batch[0])
        loss = criterion(yhat.squeeze(), batch[1].squeeze())
        total_loss += loss.item()
        pred = torch.max(yhat.data, 1)[1]
        sample_size = batch[1].size(0)
        correct_pred += (pred == batch[1]).sum().item()
        loss.backward()
        optimizer.step()
        if (i + 1) % 50 == 0:
            writer.add_scalar('Loss on Minibatches', loss.item(), step + 1)
            writer.add_scalar('Accuracy on minibatches', ((pred == batch[1]).sum().item() / sample_size), step + 1)
            step += 1
    accuracy = (correct_pred / len(CifarTrain)) * 100
    writer.add_scalar('Train Loss', np.sqrt(total_loss / len(CifarTrain)), epoch + 1)
    writer.add_scalar('Train accuracy', accuracy, epoch + 1)

    with torch.no_grad():
        test_pred = 0
        test_loss = 0
        for i, (img, label) in enumerate(test_loader):
            optimizer.zero_grad()
            output = ModelD(img)
            loss = criterion(output.squeeze(), label.squeeze())
            test_loss += loss.item()
            pred = torch.max(output.data, 1)[1]
            test_pred += (pred == label).sum().item()
        accuracy = (test_pred / len(CifarTest)) * 100
        writer.add_scalar('Test Loss', np.sqrt(test_loss / len(CifarTest)), epoch + 1)
        writer.add_scalar('Test accuracy', accuracy, epoch + 1)

writer.close()
writer.flush()


```


Loss on Minibatches



	Name	Smoothed	Value	Step	Time	Relative
Test Loss	Cifar\Dropout	0.8034	0.8034	70	Sun Jun 26, 17:53:56	3m 48s

Accuracy on minibatches



	Name	Smoothed	Value	Step	Time	Relative
Loss on minibatches	Cifar\Dropout	69.53	69.53	68	Sun Jun 26, 17:53:51	3m 43s


Training and Test Loss

Comparison with baseline

Since the probabilities chosen here are random, the efficiency can be increased more by adjusting them. For the case we can still see the decrease in the


more by adjusting them. For the case, we can still see the decrease in the training and test loss as compared to the baseline.

Train Loss



	Name	Smoothed	Value	Step	Time	Relative
●	Cifar\Dropout	0.08806	0.08806	10	Sun Jun 26, 17:53:58	3m 32s
●	Cifar\Network	0.0924	0.0924	10	Sun Jun 26, 17:10:40	1m 52s

Test Loss



	Name	Smoothed	Value	Step	Time	Relative
●	Cifar\Dropout	0.1323	0.1323	10	Sun Jun 26, 17:54:01	3m 32s
●	Cifar\Network	0.1387	0.1387	10	Sun Jun 26, 17:10:42	1m 52s



Training and Test accuracy

Comparison with baseline

The accuracy is though not as good as the one after Aigmentation+Normalization but it is better than the baseline.

Train Loss

	Name	Smoothed	Value	Step	Time	Relative
●	Cifar\Dropout	65.03	65.03	10	Sun Jun 26, 17:53:58	3m 32s
●	Cifar\Network	61.23	61.23	10	Sun Jun 26, 17:10:40	1m 52s



☐ ☐ ☐

Train Loss

	Name	Smoothed	Value	Step	Time	Relative
●	Cifar\Dropout	61.04	61.04	10	Sun Jun 26, 17:54:01	3m 32s
●	Cifar\Network	55.9	55.9	10	Sun Jun 26, 17:10:42	1m 52s

L1 Regularization

After the first fully connected layer, L1 regularization is applied to the parameters obtained from that layer. Since this layer has the maximum input among other FC layers and thus it is chosen randomly. The results may be better by using it on other or all layers. For the experiment part I have chosen this particular layer.

In []:

```
lambda1, lambda2 = 0.5, 0.5
```

In []:

```
class CNNL1(Module):
    def __init__(self, input_channels):
        super(CNNL1, self).__init__()
        self.conv1 = nn.Conv2d(input_channels, 10, kernel_size = 5)
        self.pool1 = nn.MaxPool2d(kernel_size = 2, stride = 2)
        self.conv2 = nn.Conv2d(10, 30, kernel_size = 5)
        self.conv3 = nn.Conv2d(30, 60, kernel_size = 5)
        self.pool2 = nn.MaxPool2d(kernel_size = 2, stride = 2)
        self.fc1 = nn.Linear(540, 270)
        self.fc2 = nn.Linear(270, 130)
        self.fc3 = nn.Linear(130, 10)

    def forward(self, x):
        y = F.relu(self.conv1(x))
        y = self.pool1(y)
        y = F.relu(self.conv2(y))
        y = F.relu(self.conv3(y))
        y = self.pool2(y)
        y = y.view(-1, 60*3*3)
        y1 = F.relu(self.fc1(y))
        y2 = F.relu(self.fc2(y1))
        out = self.fc3(y2)
        return out, y1
```

In []:

```


Modell1 = CNNL1(3)
criterion = nn.CrossEntropyLoss()
optimizer = torch.optim.Adam(Modell1.parameters(), lr = learning_rate)
writer = SummaryWriter(f'runs/Cifar/L1')
step = 0
for epoch in range(epochs):
    total_loss = 0
    correct_pred = 0
    sample_size = 0
    for i, batch in enumerate(train_loader):
        optimizer.zero_grad()
        yhat, y1 = Modell1(batch[0])
        entropyloss = criterion(yhat.squeeze(), batch[1].squeeze())
        y1_params = torch.cat([x.view(-1) for x in Modell1.fc1.parameters()])
        l1_regularization = lambda1*torch.norm(y1_params, 1)
        loss = entropyloss + l1_regularization
        total_loss += loss.item()
        pred = torch.max(yhat.data, 1)[1]
        sample_size = batch[1].size(0)
        correct_pred += (pred == batch[1]).sum().item()
        loss.backward()
        optimizer.step()
        if (i + 1) % 50 == 0:
            writer.add_scalar('Loss on Minibatches', loss.item(), step + 1)
            writer.add_scalar('Accuracy on minibatches', ((pred == batch[1]).sum().item()) / sample_size, step + 1)
            step += 1
    accuracy = (correct_pred / len(CifarTrain)) * 100
    writer.add_scalar('Train Loss', np.sqrt(total_loss / len(CifarTrain)), epoch + 1)
    writer.add_scalar('Train accuracy', accuracy, epoch + 1)

    with torch.no_grad():
        test_pred = 0
        test_loss = 0
        for i, (img, label) in enumerate(test_loader):
            optimizer.zero_grad()
            output, _ = Modell1(img)
            loss = criterion(output.squeeze(), label.squeeze())
            test_loss += loss.item()
            pred = torch.max(output.data, 1)[1]
            test_pred += (pred == label).sum().item()
        accuracy = (test_pred / len(CifarTest)) * 100
        writer.add_scalar('Test Loss', np.sqrt(test_loss / len(CifarTest)), epoch + 1)
        writer.add_scalar('Test accuracy', accuracy, epoch + 1)

writer.close()
writer.flush()


```

Loss on Minibatches



	Name	Smoothed	Value	Step	Time	Relative
Test Loss	Cifar\L1	10.1	10.1	70	Sun Jun 26, 14:02:45	3m 51s

Accuracy on minibatches



	Name	Smoothed	Value	Step	Time	Relative
Loss on minibatches	Cifar\L1	32.81	32.81	70	Sun Jun 26, 14:02:45	3m 51s


Training and Test Loss

comparison with baseline

Which technique is to be chosen for increasing model efficiency depends on so


many factors and thus can be varied from model to model. Here we can see after applying L1 regularization the training and test loss has increased as compared to the baseline. This is may be due to the fact that the model has penalized parameters more as lambda1 is 0.5 and thus the model is performing worse now. Due to lack of time, other values of lambda cannot be experimented.

Train Loss



Name	Smoothed	Value	Step	Time	Relative
Cifar\L1	0.2801	0.2801	10	Sun Jun 26, 14:02:47	3m 34s
Cifar\Network	0.0924	0.0924	10	Sun Jun 26, 17:10:40	1m 52s

Test Loss





Name	Smoothed	Value	Step	Time	Relative
Cifar\L1	0.1623	0.1623	10	Sun Jun 26, 14:02:50	3m 34s
Cifar\Network	0.1387	0.1387	10	Sun Jun 26, 17:10:42	1m 52s




Training and Test Accuracy

comparison with baseline

Train Loss

	Name	Smoothed	Value	Step	Time	Relative
●	Cifar\L1	33.72	33.72	10	Sun Jun 26, 14:02:47	3m 34s
●	Cifar\Network	61.23	61.23	10	Sun Jun 26, 17:10:40	1m 52s





	Name	Smoothed	Value	Step	Time	Relative
●	Cifar\L1	35.28	35.28	10	Sun Jun 26, 14:02:50	3m 34s
●	Cifar\Network	55.9	55.9	10	Sun Jun 26, 17:10:42	1m 52s

L2 Regularization

In []:

```
class CNNL2(Module):
    def __init__(self, input_channels):
        super(CNNL2, self).__init__()
        self.conv1 = nn.Conv2d(input_channels, 10, kernel_size = 5)
        self.pool1 = nn.MaxPool2d(kernel_size = 2, stride = 2)
        self.conv2 = nn.Conv2d(10, 30, kernel_size = 5)
        self.conv3 = nn.Conv2d(30, 60, kernel_size = 5)
        self.pool2 = nn.MaxPool2d(kernel_size = 2, stride = 2)
        self.fc1 = nn.Linear(540, 270)
        self.fc2 = nn.Linear(270, 130)
        self.fc3 = nn.Linear(130, 10)

    def forward(self, x):
        y = F.relu(self.conv1(x))
        y = self.pool1(y)
        y = F.relu(self.conv2(y))
        y = F.relu(self.conv3(y))
        y = self.pool2(y)
        y = y.view(-1, 60*3*3)
        y1 = F.relu(self.fc1(y))
        y2 = F.relu(self.fc2(y1))
        out = self.fc3(y2)
        return out, y1
```

In []:

```


ModelL2 = CNNL2(3)
criterion = nn.CrossEntropyLoss()
optimizer = torch.optim.Adam(ModelL2.parameters(), lr = learning_rate)
writer = SummaryWriter(f'runs/Cifar/L2')
step = 0
for epoch in range(epochs):
    total_loss = 0
    correct_pred = 0
    sample_size = 0
    for i, batch in enumerate(train_loader):
        optimizer.zero_grad()
        yhat, y1 = ModelL2(batch[0])
        entropyloss = criterion(yhat.squeeze(), batch[1].squeeze())
        y1_params = torch.cat([x.view(-1) for x in ModelL2.fc1.parameters()])
        l2_regularization = lambda2*torch.norm(y1_params, 2)
        loss = entropyloss + l2_regularization
        total_loss += loss.item()
        pred = torch.max(yhat.data, 1)[1]
        sample_size = batch[1].size(0)
        correct_pred += (pred == batch[1]).sum().item()
        loss.backward()
        optimizer.step()
        if (i + 1) % 50 == 0:
            writer.add_scalar('Loss on Minibatches', loss.item(), step + 1)
            writer.add_scalar('Accuracy on minibatches', ((pred == batch[1]).sum().item()) / sample_size, step + 1)
            step += 1
    accuracy = (correct_pred / len(CifarTrain)) * 100
    writer.add_scalar('Train Loss', np.sqrt(total_loss / len(CifarTrain)), epoch + 1)
    writer.add_scalar('Train accuracy', accuracy, epoch + 1)

    with torch.no_grad():
        test_pred = 0
        test_loss = 0
        for i, (img, label) in enumerate(test_loader):
            optimizer.zero_grad()
            output, _ = ModelL2(img)
            loss = criterion(output.squeeze(), label.squeeze())
            test_loss += loss.item()
            pred = torch.max(output.data, 1)[1]
            test_pred += (pred == label).sum().item()
        accuracy = (test_pred / len(CifarTest)) * 100
        writer.add_scalar('Test Loss', np.sqrt(test_loss / len(CifarTest)), epoch + 1)
        writer.add_scalar('Test accuracy', accuracy, epoch + 1)

writer.close()
writer.flush()


```

Loss on Minibatches



	Name	Smoothed	Value	Step	Time	Relative
Test Loss	Cifar\L2	1.51	1.51	70	Sun Jun 26, 14:35:06	3m 46s

Accuracy on minibatches



	Name	Smoothed	Value	Step	Time	Relative
Loss on Minibatches	Cifar\L2	52.34	52.34	70	Sun Jun 26, 14:35:06	3m 46s


Training and Test loss

Comparison with baseline

As seen above for the L1 regularization, same is observed for the L2


regularization. There has been an increase in the loss and decrease in the accuracy as compared to the baseline.

Train Loss



	Name	Smoothed	Value	Step	Time	Relative
Train accuracy	Cifar\l2	0.11	0.11	10	Sun Jun 26, 14:35:09	3m 29s
	Cifar\Network	0.0924	0.0924	10	Sun Jun 26, 17:10:40	1m 52s


Test Loss



	Name	Smoothed	Value	Step	Time	Relative
Test accuracy	Cifar\l2	0.1482	0.1482	10	Sun Jun 26, 14:35:11	3m 29s
	Cifar\Network	0.1387	0.1387	10	Sun Jun 26, 17:10:42	1m 52s

training and Test Accuracy

comparison with baseline



Optimizers

Now we will be experimenting with two optimizers and how they response the

localhost:8888/notebooks/DDALab/Exercise7/Report_DDA_Ex7.ipynb

37/47

different learning rates.

In []:

```
learningRate = [0.01, 0.001, 0.0001]
```

Adam Optimizer

In []:


```
for learning_rate in learningRate:
    Model = CNNNetwork(3)
    criterion = nn.CrossEntropyLoss()
    optimizer = torch.optim.Adam(Model.parameters(), lr = learning_rate)
    writer = SummaryWriter(f'runs/Cifar/Adam {learning_rate}')
    step = 0
    for epoch in range(epochs):
        total_loss = 0
        correct_pred = 0
        sample_size = 0
        for i, batch in enumerate(train_loader):
            optimizer.zero_grad()
            yhat = Model(batch[0])
            loss = criterion(yhat.squeeze(), batch[1].squeeze())
            total_loss += loss.item()
            pred = torch.max(yhat.data, 1)[1]
            sample_size = batch[1].size(0)
            correct_pred += (pred == batch[1]).sum().item()
            loss.backward()
            optimizer.step()
            if (i + 1) % 50 == 0:
                writer.add_scalar('Loss on Minibatches', loss.item(), step + 1)
                writer.add_scalar('Accuracy on minibatches', ((pred == batch[1]).sum().item() /
                                                                sample_size),
                                  step + 1)
            step += 1

        accuracy = (correct_pred / len(CifarTrain)) * 100
        writer.add_scalar('Train Loss', np.sqrt(total_loss / len(CifarTrain)), epoch + 1)
        writer.add_scalar('Train accuracy', accuracy, epoch + 1)

    with torch.no_grad():
        test_pred = 0
        test_loss = 0
        for i, (img, label) in enumerate(test_loader):
            optimizer.zero_grad()
            output = Model(img)
            loss = criterion(output.squeeze(), label.squeeze())
            test_loss += loss.item()
            pred = torch.max(output.data, 1)[1]
            test_pred += (pred == label).sum().item()
        accuracy = (test_pred / len(CifarTest)) * 100
        writer.add_scalar('Test Loss', np.sqrt(test_loss / len(CifarTest)), epoch + 1)
        writer.add_scalar('Test accuracy', accuracy, epoch + 1)


    writer.close()
    writer.flush()
```

Loss on Minibatches



Name	Smoothed	Value	Step	Time	Relative
Cifar\Adam 0.0001	1.554	1.554	70	Sun Jun 26, 16:33:55	3m 34s
Cifar\Adam 0.001	0.8094	0.8094	70	Sun Jun 26, 16:30:15	3m 34s
Cifar\Adam 0.01	2.307	2.307	70	Sun Jun 26, 16:26:34	2m 57s

Accuracy on minibatches




Name	Smoothed	Value	Step	Time	Relative
Cifar\Adam 0.0001	41.41	41.41	70	Sun Jun 26, 16:33:55	3m 34s
Cifar\Adam 0.001	75.78	75.78	70	Sun Jun 26, 16:30:15	3m 34s
Cifar\Adam 0.01	8.594	8.594	70	Sun Jun 26, 16:26:34	2m 57s


Training and Test Loss

As can be seen below, both the training and test loss has decreased more faster when the initial learning rate is neither too big nor too small. For example here three different learning rates were experimented and the results have been found best for learning rate 0.001.

Train Loss



	Name	Smoothed	Value	Step	Time	Relative
●	Cifar\Adam 0.0001	0.1082	0.1082	10	Sun Jun 26, 16:33:58	3m 19s
●	Cifar\Adam 0.001	0.08308	0.08308	10	Sun Jun 26, 16:30:17	3m 18s
●	Cifar\Adam 0.01	0.1342	0.1342	10	Sun Jun 26, 16:26:36	2m 45s




Training and Test Accuracy

since the learning rate that has resulted in maximum decrease in loss has also increased the accuracy most.


Train Loss




	Name	Smoothed	Value	Step	Time	Relative
●	Cifar\Adam 0.0001	45.06	45.06	10	Sun Jun 26, 16:33:58	3m 19s
●	Cifar\Adam 0.001	68.74	68.74	10	Sun Jun 26, 16:30:17	3m 18s
●	Cifar\Adam 0.01	10.01	10.01	10	Sun Jun 26, 16:26:36	2m 45s

Train accuracy



Test accuracy





	Name	Smoothed	Value	Step	Time	Relative
●	Cifar\Adam 0.0001	45.41	45.41	10	Sun Jun 26, 16:34:01	3m 20s
●	Cifar\Adam 0.001	64.26	64.26	10	Sun Jun 26, 16:30:19	3m 18s
●	Cifar\Adam 0.01	10	10	10	Sun Jun 26, 16:26:38	2m 45s

SGD Optimizer

In []:

```

for learning_rate in learningRate:
    Model = CNNNetwork(3)
    criterion = nn.CrossEntropyLoss()
    optimizer = torch.optim.SGD(Model.parameters(), lr = learning_rate)
    writer = SummaryWriter(f'runs/Cifar/SGD {+ learning_rate}')
    step = 0
    for epoch in range(epochs):
        total_loss = 0
        correct_pred = 0
        sample_size = 0
        for i, batch in enumerate(train_loader):
            optimizer.zero_grad()
            yhat = Model(batch[0])
            loss = criterion(yhat.squeeze(), batch[1].squeeze())
            total_loss += loss.item()
            pred = torch.max(yhat.data, 1)[1]
            sample_size = batch[1].size(0)
            correct_pred += (pred == batch[1]).sum().item()
            loss.backward()
            optimizer.step()
            if (i + 1) % 50 == 0:
                writer.add_scalar('Loss on Minibatches', loss.item(), step + 1)
                writer.add_scalar('Accuracy on minibatches', ((pred == batch[1]).sum().item() /
                                                                sample_size),
                                   step + 1)
                step += 1


        accuracy = (correct_pred / len(CifarTrain)) * 100
        writer.add_scalar('Train Loss', np.sqrt(total_loss / len(CifarTrain)), epoch + 1)
        writer.add_scalar('Train accuracy', accuracy, epoch + 1)

    with torch.no_grad():
        test_pred = 0
        test_loss = 0
        for i, (img, label) in enumerate(test_loader):
            optimizer.zero_grad()
            output = Model(img)
            loss = criterion(output.squeeze(), label.squeeze())
            test_loss += loss.item()
            pred = torch.max(output.data, 1)[1]
            test_pred += (pred == label).sum().item()
        accuracy = (test_pred / len(CifarTest)) * 100
        writer.add_scalar('Test Loss', np.sqrt(test_loss / len(CifarTest)), epoch + 1)
        writer.add_scalar('Test accuracy', accuracy, epoch + 1)

writer.close()
writer.flush()


```

Loss on Minibatches



	Name	Smoothed	Value	Step	Time	Relative
Test Loss	Cifar\SGD 0.0001	2.302	2.302	64	Sun Jun 26, 16:51:04	3m 8s
	Cifar\SGD 0.001	2.298	2.298	64	Sun Jun 26, 16:47:35	3m 6s
Test Accuracy	Cifar\SGD 0.01	2.048	2.048	64	Sun Jun 26, 16:44:07	2m 42s

Accuracy on minibatches




	Name	Smoothed	Value	Step	Time	Relative
Loss on Minibatches	Cifar\SGD 0.0001	14.06	14.06	68	Sun Jun 26, 16:51:14	3m 18s
	Cifar\SGD 0.001	10.94	10.94	68	Sun Jun 26, 16:47:44	3m 16s
Test Loss	Cifar\SGD 0.01	22.66	22.66	68	Sun Jun 26, 16:44:16	2m 52s

Training and Test Loss


SGD seems to work better with not too small initial learning rate and afterwards a small learning rate there is hardly any much declination in the loss.

Train Loss



	Name	Smoothed	Value	Step	Time	Relative
Train accuracy	Cifar\SGD 0.0001	0.1342	0.1342	8	Sun Jun 26, 16:50:39	2m 26s
	Cifar\SGD 0.001	0.1342	0.1342	8	Sun Jun 26, 16:47:10	2m 25s
	Cifar\SGD 0.01	0.1272	0.1272	8	Sun Jun 26, 16:43:44	2m 5s

Test Loss



	Name	Smoothed	Value	Step	Time	Relative
Test accuracy	Cifar\SGD 0.0001	0.1902	0.1902	8	Sun Jun 26, 16:50:41	2m 26s
	Cifar\SGD 0.001	0.1901	0.1901	8	Sun Jun 26, 16:47:12	2m 25s
Train Loss	Cifar\SGD 0.01	0.1787	0.1787	8	Sun Jun 26, 16:43:46	2m 5s


Training and Test Accuracy

As seen above for loss same applies to accuracy.


Train Loss

	Name	Smoothed	Value	Step	Time	Relative
●	Cifar\SGD 0.0001	10	10	8	Sun Jun 26, 16:50:39	2m 26s
●	Cifar\SGD 0.001	10	10	8	Sun Jun 26, 16:47:10	2m 25s
●	Cifar\SGD 0.01	22.23	22.23	8	Sun Jun 26, 16:43:44	2m 5s

Train accuracy



Test accuracy



	Name	Smoothed	Value	Step	Time	Relative
Train Loss	Cifar\SGD 0.0001	10.02	10.02	9	Sun Jun 26, 16:51:02	2m 47s
Train Loss	Cifar\SGD 0.001	10	10	9	Sun Jun 26, 16:47:32	2m 46s
Train accuracy	Cifar\SGD 0.01	24.13	24.13	9	Sun Jun 26, 16:44:05	2m 24s

In []: